

# SparseDR: Differentiable Rendering of Sparse Signed Distance Fields

First Author<sup>1</sup>, Second Author<sup>2</sup>, Third Author<sup>2,3</sup> and Fourth Author<sup>4</sup>

<sup>1</sup>First Affiliation

<sup>2</sup>Second Affiliation

<sup>3</sup>Third Affiliation

<sup>4</sup>Fourth Affiliation

{first, second}@example.com, third@other.example.com, fourth@example.com

## Abstract

We present SparseDR, a differentiable rendering implementation designed for sparse representations based on Signed Distance Fields (SDF). We leverage the Sparse Brick Set (SBS) representation and propose an adaptation of redistancing and regularization of the SDF defined on SBS. This enables our method to surpass existing differentiable rendering implementations in accuracy at limited memory budgets by increasing the effective resolution of the SDF representation.

## 1 Introduction

Signed distance functions (SDFs) provide a continuous, implicit representation of geometry, defining the shortest Euclidean distance from any point in space to the nearest surface, with the sign indicating interior or exterior regions. This enables gradient-based optimization by supporting differentiable rendering, where ray marching samples distances to compute density and color, facilitating inverse rendering tasks such as shape and material recovery from images.

SDF-based differentiable rendering produces high-fidelity reconstructions with smooth surfaces and topological consistency, outperforming many existing methods in handling complex topologies. Unfortunately, existing SDF-based differentiable rendering implementations are constrained in reconstruction accuracy due to high memory consumption and substantial computational demands. Our work proposes a solution to this problem, alleviating limitations through better data structures and sampling techniques:

- We propose an adaptation of SDF-based differentiable rendering to a sparse data structure, which increases the effective resolution of the SDF representation and thereby improves the surface reconstruction accuracy.
- We provide two novel algorithms of redistancing: a full version for SBS in general and its accelerated version specifically for the differentiable rendering of SBS.

## 2 Related Work

[Vicini *et al.*, 2022] introduces differentiable rendering operating directly on an SDF representation. The authors modi-

fied sphere tracing so that, in addition to computing the surface intersection, it evaluates a special reparameterization of rays, enabling correct gradients with respect to shape parameters, including the visibility contribution of individual surface regions. This work also first proposed the use of a *redistancing* — a process of recomputing values on a grid such that all grid cells store correct distances to the surface. In differentiable rendering this recomputation is required because, after an optimization step and updating the distances at grid nodes, the grid no longer strictly represents an SDF. [Wang *et al.*, 2024] exploits properties of distance functions and estimate the boundary integral by an integral over a relaxed boundary, i.e., a thin band around the object silhouette. [Zhang *et al.*, 2025] propose for computing gradients in the presence of multiple surfaces along a ray path: instead of locally differentiating a single surface, the authors consider a volumetric, non-local perturbation of the surface.

All of the above methods represent the distance function volumetrically on a regular grid, which leads to three fundamental issues: (1) significant memory consumption, especially when using automatic differentiation frameworks such as DrJit [Jakob *et al.*, 2022] or PyTorch [Jason and Yang, 2024]; (2) slow ray–surface intersection based on sphere tracing; (3) an expensive redistancing algorithm that propagates modified SDF values across the entire grid.

Hash-table-based methods such as InstantNGP [Müller *et al.*, 2022] could mitigate the first issue, but by themselves cannot address the second and third. The key difficulty is that the data structure used to store the SDF directly constrains how differentiable ray–surface intersections are computed and how redistancing is performed, and implementing redistancing on hash tables is, in our view, highly nontrivial. Our work shows how these three components – data structure, ray intersection, and redistancing – can be joined together in an efficient manner.

## 3 System Implementation

To store the surface, SparseDR uses the Sparse Brick Set (SBS) proposed in [Söderlund *et al.*, 2022], which is a hybrid of hierarchical and regular representations, see Figure 1. The entire scene is encoded as a BVH or an octree, while its leaves store small regular grids, for example 4x4x4. For gradient computation, we adapt the relaxed boundary method from [Wang *et al.*, 2024], which we refer to as our baseline. At the

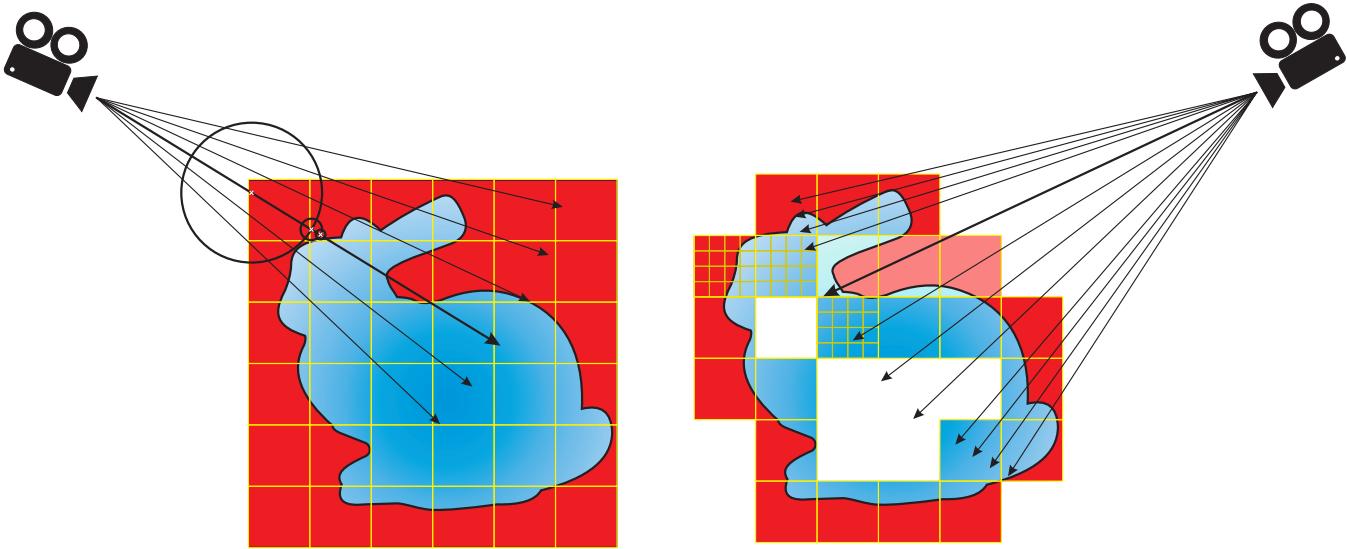


Figure 1: Comparison of SparseDR algorithm (on the right) with the baseline. The key differences are using sparse SDF representation, SBS, and intersecting voxels along the ray instead of sphere tracing.

beginning of the optimization process, the distance function is initialized with a sphere defined on a low-resolution grid (32x32x32). At each iteration, we estimate the gradient and update the distance values using a gradient descent step, followed by a regularization and accelerated redistancing (see Figure 2). On every iteration when upsampling is required, which is defined by the user, three steps are executed: sparsification step, full redistancing and upsampling. Sparsification and full redistancing are also applied before the resulting model is saved.

### 3.1 Gradient evaluation

The key distinction of the method underlying SparseDR lies in the transition from an SDF grid to a sparse brick set. In our case, each brick has a fixed size of 4x4x4 voxels. For the SBS, a BVH tree is constructed, with each leaf node storing a single brick. In addition to sphere tracing, alternative ray–SDF intersection methods have been implemented: Newton and analytic, specifically designed to find intersections inside the voxel. A ray first traverses the BVH; if an intersection with a leaf node is found, it then intersects the voxels of the brick. The ray proceeds through each brick along its path until the first intersection with a surface is reached.

The alternative ray–SDF intersection methods also support searching for relaxed boundary points. According to [reference], a point belongs to the relaxed boundary if the SDF value at that point is below a certain threshold called the relaxed epsilon, and the SDF derivative along the ray reaches its local minimum. The Newton and analytic methods are based on the fact that inside a voxel, the SDF along the ray can be represented as a cubic polynomial. Thus, within the voxel, a potential relaxed boundary point can be found as the root of a quadratic equation, after which it is sufficient to verify that the SDF value at this point is less than the relaxed epsilon. The case when the local minimum occurs on the boundary between voxels or bricks is handled separately.

SparseDR is implemented in C++ and Vulkan and does not use automatic differentiation. Gradients are accumulated in a float array whose size equals the number of updated distances. In the shader, when computing the color of a pixel for 16 samples, information about the voxels hit by the samples is stored; after every 16 samples (or after the last one), the gradients are added to the array. SparseDR optionally uses an additional shader that samples only pixels containing boundaries and evaluates only the boundary integral, without the need for color computation. This adaptive boundary sampling allows cheaper samples and yields more accurate gradients for smaller relaxed epsilon values, thereby reducing bias.

### 3.2 Redistancing

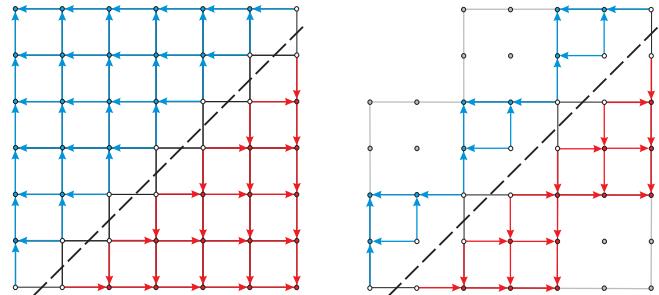


Figure 2: Comparison of the redistancing algorithm (on the left) with the accelerated version. Updating only the bricks with a surface is sufficient to get correct relaxed boundary points, as the brick size is larger than the relaxed epsilon.

For differentiable rendering of SBS in SparseDR, two versions of redistancing have been implemented. The first, full version, is an adaptation of the parallel fast sweeping method for SBS, which is used to acquire the correct SDF. In this version, it was necessary to address the problem of distance

135 updates across empty space between bricks. Consider a brick  
136 without neighbors and without any surface inside it. Following  
137 the redistancing algorithm, the distances within it will be  
138 initialized with a constant value and will never be updated.  
139 In differentiable rendering, such extreme cases do not occur,  
140 but situations may arise where the distances within a brick  
141 have larger absolute values than they should, since adjacent  
142 bricks are not taken into account. To solve this issue, an ex-  
143 trapolation step between brick faces was introduced. For each  
144 face of a brick without a direct neighbor, the nearest neighbor  
145 bricks are found, and one of their faces is selected for extrap-  
146 olating the values. These neighbors are determined once, and  
147 after each fast sweeping step, values are extrapolated consid-  
148 ering the grid spacing. From the extrapolated and previous  
149 values, the one with the smallest absolute value is chosen. To  
150 propagate the updated distances, the fast sweeping method is  
151 always applied as the final step of this algorithm.

152 However, using SBS instead of a grid and Newton’s  
153 method for intersection search inside the voxel instead of  
154 sphere tracing opened opportunities for optimization of the  
155 redistancing procedure. These resulted in an accelerated ver-  
156 sion, which is used after each iteration. The first idea is to  
157 exploit the properties of intersection based on the Newton  
158 method. Since each ray traverses all voxels along its path, it  
159 does not depend on exact distance values in voxels that do not  
160 contain boundaries. Therefore, to obtain accurate surface in-  
161 tersections and relaxed boundary points, it is sufficient to ap-  
162 ply redistancing only for surface bricks. The second idea as-  
163 sumes that in surface bricks it is enough to compute distances  
164 considering only the surface stored within the same brick,  
165 without propagating distance values between bricks. This as-  
166 sumption was only partially confirmed, since during the ini-  
167 tialization stage, redistancing recalculates distances defining  
168 the surface, and they must be as accurate as possible. How-  
169 ever, after initialization, redistancing localized within indi-  
170 vidual bricks is sufficient. Figure Nn shows the MSE plot  
171 for optimization iterations of the Nefertiti bust model using  
172 both redistancing versions. The accelerated version neither  
173 degrades reconstruction quality nor affects convergence to the  
174 reference.

### 175 3.3 Sparsification and Upsampling

176 The sparsification step is performed before upsampling. Dur-  
177 ing this stage, all bricks containing the surface are preserved,  
178 as well as all bricks within a radius of 1 voxel from them,  
179 including diagonal neighbors. If any of the adjacent bricks  
180 are missing, they are added at this stage. This is necessary  
181 in cases where the reconstructed object contains fine details,  
182 so that the bricks where these details’ reconstruction should  
183 reside would otherwise be absent. Thus, this step not only  
184 contributes to reducing the model size, but also eliminates a  
185 potential limitation compared to the baseline method. After  
186 the sparsification step, the full version of redistancing is ap-  
187 plied to recompute all distances before upsampling or saving  
188 the resulting SBS, ensuring that the model represents a valid  
189 SDF. The upsampling step doubles the resolution of the re-  
190 maining bricks, generating eight new bricks from each origi-  
191 nal one. The new distance values are computed using trilinear  
192 interpolation.

## 4 Experiments

193 Experiments with the baseline method were conducted  
194 using [https://github.com/zichenwang01/relaxed-boundary/],  
195 which provides its official implementation. The baseline  
196 method employed the hyperparameters specified in the pa-  
197 per [reference], while the remaining parameters were taken  
198 from the configuration file “turbo.json”. The only modifi-  
199 cations concerned the number of viewpoints, the batch size,  
200 and the steps at which upsampling was performed. For con-  
201 sistency across experiments, an identical lighting setup was  
202 used: two directional light sources with directions  $(1, 1, 1)$   
203 and  $(-1, -1, -1)$ , along with ambient light. The experiments  
204 were carried out on six models. Reconstruction used 16 view-  
205 points uniformly distributed over a sphere around the scene,  
206 generated by the algorithm from [reference]. The batch size is  
207 4 viewpoints, image resolution is  $1024 \times 1024$ . The full opti-  
208 mization process comprised 1000 steps, with upsampling ap-  
209 plied at steps 200, 400, 600, 700, 800, and 900. Computa-  
210 tions were performed on an AMD Ryzen 9 7950X CPU and  
211 an NVIDIA GeForce RTX 4090 GPU.

## 5 Conclusion

213 In this work, we present SparseDR, a system for 3D recon-  
214 struction based on the method of differentiable rendering for  
215 a sparse SDF representation. Using [reference] as a base-  
216 line, we leverage the advantages of SBS and modify all stages  
217 of the algorithm. As a result, SparseDR achieves perfor-  
218 mance and reconstruction quality comparable to the original  
219 approach while reducing memory consumption by an order of  
220 magnitude. Future work will focus on upgrading SparseDR  
221 to a full-fledged path tracer and addressing the challenges that  
222 complicate the practical use of differentiable rendering, in  
223 particular the need for precise camera parameter information.

### 225 5.1 Order of Sections

226 Sections should be arranged in the following order [?]:

- 227 1. Main content sections (numbered)
- 228 2. Appendices (optional, numbered using capital letters)
- 229 3. Ethical statement (optional, unnumbered)
- 230 4. Acknowledgements (optional, unnumbered)
- 231 5. Contribution statement (optional, unnumbered)
- 232 6. References (required, unnumbered)

## 233 References

- [Jakob *et al.*, 2022] Wenzel Jakob, Sébastien Speierer, Nico-  
234 las Roussel, and Delio Vicini. Drjit: a just-in-time com-  
235 piler for differentiable rendering. 41(4), 2022.
- [Jason and Yang, 2024] Jason and Edward He *et al.* Yang.  
236 Pytorch 2: Faster machine learning through dynamic  
237 python bytecode transformation and graph compilation. In  
238 *Proceedings of the 29th ACM International Conference  
239 on Architectural Support for Programming Languages  
240 and Operating Systems, Volume 2, ASPLOS ’24, page  
241 929–947, New York, NY, USA, 2024. Association for  
242 Computing Machinery.*

- 245 [Müller *et al.*, 2022] Thomas Müller, Alex Evans, Christoph  
246 Schied, and Alexander Keller. Instant neural graphics  
247 primitives with a multiresolution hash encoding. *ACM*  
248 *Trans. Graph.*, 41(4), 2022.
- 249 [Söderlund *et al.*, 2022] Herman Hansson Söderlund, Alex  
250 Evans, and Tomas Akenine-Möller. Ray tracing of signed  
251 distance function grids. *Journal of Computer Graphics*  
252 *Techniques Vol*, 11(3), 2022.
- 253 [Vicini *et al.*, 2022] Delio Vicini, Sébastien Speierer, and  
254 Wenzel Jakob. Differentiable signed distance function ren-  
255 dering. *ACM Trans. Graph.*, 41(4), July 2022.
- 256 [Wang *et al.*, 2024] Zichen Wang, Xi Deng, Ziyi Zhang,  
257 Wenzel Jakob, and Steve Marschner. A simple approach to  
258 differentiable rendering of sdfs. In *SIGGRAPH Asia 2024*  
259 *Conference Papers*, New York, NY, USA, 2024. Associa-  
260 tion for Computing Machinery.
- 261 [Zhang *et al.*, 2025] Ziyi Zhang, Nicolas Roussel, and Wen-  
262 zel Jakob. Many-worlds inverse rendering. *ACM Trans.*  
263 *Graph.*, 45(1), October 2025.